

## Importing important libraries

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

## Importing data

```
In [3]: data = pd.read_csv("Titanic-Dataset.csv")
```

In [4]:

data

Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fa
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05
...	...	...	...	...	...	...	...	...	...	...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.00
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.00
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.45
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.00
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.75

891 rows × 12 columns



In [10]:

new\_data = data.drop(['PassengerId'], axis=1)

```
In [11]: new_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   Survived      891 non-null    int64  
 1   Pclass        891 non-null    int64  
 2   Name          891 non-null    object  
 3   Sex           891 non-null    object  
 4   Age           714 non-null    float64 
 5   SibSp         891 non-null    int64  
 6   Parch         891 non-null    int64  
 7   Ticket        891 non-null    object  
 8   Fare          891 non-null    float64 
 9   Cabin         204 non-null    object  
10   Embarked      889 non-null    object  
dtypes: float64(2), int64(4), object(5)
memory usage: 76.7+ KB
```

```
In [12]: new_data.isnull().sum()
```

```
Out[12]: Survived      0
Pclass      0
Name        0
Sex         0
Age         177
SibSp       0
Parch       0
Ticket      0
Fare        0
Cabin       687
Embarked    2
dtype: int64
```

## Handling missing values

```
In [13]: titanic_data = new_data.drop(['Cabin'], axis =1)
```

In [14]: `titanic_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    891 non-null    int64
 1   Pclass      891 non-null    int64
 2   Name        891 non-null    object
 3   Sex         891 non-null    object
 4   Age         714 non-null    float64
 5   SibSp       891 non-null    int64
 6   Parch       891 non-null    int64
 7   Ticket      891 non-null    object
 8   Fare        891 non-null    float64
 9   Embarked    889 non-null    object
dtypes: float64(2), int64(4), object(4)
memory usage: 69.7+ KB
```

In [15]: `titanic_data['Age'].fillna(titanic_data['Age'].mean(),inplace=True)`

In [16]: `titanic_data.isnull().sum()`

```
Out[16]: Survived      0
Pclass      0
Name        0
Sex         0
Age         0
SibSp       0
Parch       0
Ticket      0
Fare        0
Embarked    2
dtype: int64
```

In [17]: *# Finding the mode value of Embarked column.*  
`print(titanic_data['Embarked'].mode())`

```
0    S
Name: Embarked, dtype: object
```

In [21]: *# Replacing the missing values in embarked column with mode value.*  
`titanic_data['Embarked'].fillna(titanic_data['Embarked'].mode()[0], inplace`

In [2]: `titanic_data.isnull().sum()`

```
-----
-
NameError                                Traceback (most recent call las
t)
Cell In[2], line 1
----> 1 titanic_data.isnull().sum()

NameError: name 'titanic_data' is not defined
```

## Data Analysis

In [23]:

titanic\_data.describe()

Out[23]:

	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	13.002015	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	22.000000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	29.699118	0.000000	0.000000	14.454200
75%	1.000000	3.000000	35.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [24]:

titanic\_data['Survived'].value\_counts()

Out[24]:

0 549  
1 342  
Name: Survived, dtype: int64


## Data Visualization

In [29]:

*# Count plot - survived column*  
sns.countplot(x='Survived', data= titanic\_data)


Out[29]:

<Axes: xlabel='Survived', ylabel='count'>




```
In [30]: sns.countplot(x='Sex', data= titanic_data)
```

```
Out[30]: <Axes: xlabel='Sex', ylabel='count'>
```




```
In [32]: sns.countplot(x="Sex",hue="Survived",data=titanic_data)
```

```
Out[32]: <Axes: xlabel='Sex', ylabel='count'>
```



```
In [33]: sns.countplot(x="Pclass",hue="Survived",data=titanic_data)
```

```
Out[33]: <Axes: xlabel='Pclass', ylabel='count'>
```



## Encoding the categorical columns

```
In [36]: titanic_data.replace({'Sex': {'male': 0, 'female': 1}, 'Embarked': {'S': 0,
```

```
In [37]: titanic_data.head()
```

```
Out[37]:
```

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2500	0
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1	0	PC 17599	71.2833	1
2	1	3	Heikkinen, Miss. Laina	1	26.0	0	0	STON/O2. 3101282	7.9250	0
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	35.0	1	0	113803	53.1000	0
4	0	3	Allen, Mr. William Henry	0	35.0	0	0	373450	8.0500	0



## Separating features and targets

```
In [42]: x = titanic_data.drop(["Name", "Ticket", "Survived"], axis=1)
```

```
In [43]: y = titanic_data['Survived']
```

```
In [44]: print(x)
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	0	22.000000	1	0	7.2500	0
1	1	1	38.000000	1	0	71.2833	1
2	3	1	26.000000	0	0	7.9250	0
3	1	1	35.000000	1	0	53.1000	0
4	3	0	35.000000	0	0	8.0500	0
..	...	...	...	...	...	...	...
886	2	0	27.000000	0	0	13.0000	0
887	1	1	19.000000	0	0	30.0000	0
888	3	1	29.699118	1	2	23.4500	0
889	1	0	26.000000	0	0	30.0000	1
890	3	0	32.000000	0	0	7.7500	2

[891 rows x 7 columns]

```
In [45]: print(y)
```

```
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
```

Name: Survived, Length: 891, dtype: int64

## Splitting the data into training and testing data

```
In [46]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_s
```

```
In [48]: print(x.shape, x_train.shape, x_test.shape)
```

(891, 7) (712, 7) (179, 7)

## Training model

```
In [49]: model = LogisticRegression()
```

```
In [50]: model.fit(x_train,y_train)
```

C:\Users\Apoorva\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
Out[50]: LogisticRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

## Model evaluation

```
In [51]: x_train_prediction = model.predict(x_train)
```

```
x_train_prediction
```

```
In [53]: training_data_accuracy = accuracy_score(y_train,x_train_prediction)
```

```
In [54]: print("Accuracy score of trainig data:",training_data_accuracy )
```

```
Accuracy score of trainig data: 0.8075842696629213
```

```
In [55]: x_test_prediction = model.predict(x_test)
```

```
In [56]: print(x_test_prediction)
```

```
[0 0 1 0 0 0 0 0 0 0 0 1 1 0 0 1 0 0 1 0 1 1 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1
 0 0 0 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 1 0 1 0
 1 0 0 0 1 0 1 0 0 0 1 1 0 0 1 0 0 0 0 0 0 1 0 1 0 0 1 0 1 1 0 1 1 0 0 0 0
 0 0 0 1 1 0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 1 1 1 1 0 1 0 0
 0 1 0 0 0 0 1 0 0 1 1 0 1 0 0 0 1 1 0 0 1 0 0 1 1 1 0 0 0 0 0 0]
```

```
In [58]: test_data_accuracy = accuracy_score(y_test,x_test_prediction)
```

```
In [60]: print("Accuracy score of test data", test_data_accuracy)
```

Accuracy score of test data 0.7821229050279329

```
In [ ]:
```